

Final Report — Human Activity Recognition on 1 Hz Wrist Accelerometry

NYCU 535703 Data Mining (Spring 2026), Assignment 3 · Kaggle team 314540066

0. Basics (code, how to run, reproducibility)

- **Public GitHub repository:** <https://github.com/Ashurali/DM2026-Assignment-3-MKS>
- **Kaggle competition:** `nycu-data-mining-assignment-3` · best **public LB = 0.8234** (submissions/sub_pc_b20.csv).

- **One-command reproduction (no training, < 1 min):**

```
python -m venv .venv
# Windows: .\.venv\Scripts\Activate.ps1 | Linux/Mac: source .venv/bin/activate
pip install -r requirements.txt
python scripts/reproduce_final.py # -> submissions/sub_pc_b20.csv (== public LB)
```

The frozen base-model out-of-fold (OOF) + test probabilities are committed in `oof/`, so the winning submission regenerates **deterministically from a clean clone**. `reproduce_final.py` runs the full final pipeline and asserts the exact winning class counts (L2 = 314, L3 = 559). A notebook version is in `notebooks/03_reproduce_final.ipynb`.

- **Full pipeline (retrain from raw data):** see `README.md` → *Full pipeline* (Phases A-E).
 - **Reproducibility statement:** report, code, and the Kaggle score are consistent — `reproduce_final.py` emits the exact CSV that scored 0.8234, deterministically, from artifacts committed in the repo.
-

1. Problem & data

Six-class human-activity recognition; metric = **macro-F1** (every class weighs equally). The signal is a **1 Hz aggregated tri-axial accelerometer**: each file holds 300 seconds × six columns `mean_x`, `mean_y`, `mean_z`, `std_x`, `std_y`, `std_z` (per-second mean and standard deviation), with **one activity label per file**. The data is the Bruno et al. (2014) wrist-worn ADL set: **gravity is not removed**, the device is **wrist-mounted**, and there is **no gyroscope**.

	Files	Users	rows/file	NaNs
Train	11,020	60 (User_001-060)	300	0
Test	6,849	40 (User_061-100)	300	0

Train and test users are completely disjoint — the defining difficulty (cross-subject generalization).

2. Preliminary analysis (Q1)

Five observations from the raw data drove every design decision.

1. **Data is pristine.** Every file is exactly 300 rows, zero NaNs, no duplicate `file_id`. No imputation needed.

2. **Severe class imbalance — 33x.** Counts {L0:4643, L1:4695, L2:358, L3:656, L4:142, L5:526}; L0+L1 = 84.7%. For macro-F1, the rare classes dominate the difficulty.
3. **Disjoint train/test users.** Overlap = 0 → **GroupKFold(5, groups=user_id) is mandatory everywhere**, user_id is forbidden as a feature, and any normalization must be **per-file**.
4. **An obvious intensity gradient, and one hard cluster.** Per-class mean of std_* runs monotonically from ≈ 0.009 (L0, near-static) to ≈ 0.36 (L4, high-intensity). On a t-SNE of the six raw features, **L4 forms a distinct cluster while L0/L1/L2/L3/L5 overlap heavily** — and **L2 (only 358 files) is the bottleneck**, buried inside the dominant L1 walking class.
5. **1 Hz aliases gait.** True gait (~ 2 Hz) is above the 0.5 Hz Nyquist of the 1 Hz series, but the std_* channels preserve within-second high-frequency energy implicitly.

Naive baselines (GroupKFold-5 CV macro-F1):

Model	Features	CV macro-F1
Majority class	—	0.1331
Logistic regression (balanced)	6 column means	0.5315
LightGBM	271-feature catalog (see §3)	0.7091

Design implications: (a) attack the *overlapping low-intensity classes* with richer features and sequence structure, not just the easy L4; (b) make everything *cross-subject-robust* (group CV, no user-identifying features); (c) recover *minority-class recall* lost to the softmax bias.

3. Preprocessing & feature engineering (Q2)

No cleaning was required (§2.1); preprocessing is therefore **feature construction + calibration**. Each technique below is reported with its measured effect (GroupKFold-5 OOF macro-F1 unless noted).

Technique	What it adds	Effect
Engineered feature catalog (271 features: 11 families — basic-stats, FFT, autocorr, sub-window pooling, jerk, zero-cross, magnitude, gravity, cross-axis, per-file-norm, quality)	turns 6 raw means into rich per-channel statistics + time-domain descriptors	LR 0.5315 → LGBM 0.7091 (+0.176)
catch22 + extended families + base-model stacking (combo_full_v2, an 805-column LightGBM stacker over base-model OOFs + features)	canonical time-series features + ensemble meta-features	0.7091 → ≈ 0.79 OOF
Hierarchical pipeline + α-blend (P1 flat stacker + P2 coarse→fine hierarchy, blended 0.842/0.158)	decouples easy boundaries from the hard L1↔L2 sub-problem	→ 0.7880 OOF / 0.8154 LB
Per-class isotonic calibration (5-fold GroupKFold OOF)	corrects softmax's majority-class bias → recovers minority recall	folded into the 0.7880 / 0.8154 result; prerequisite for the threshold grid
Per-file z-score normalization	removes per-user wrist-orientation offset	redundant here (ablation $\Delta = -0.001$, within noise) — tree models absorb the drift internally

The feature catalog is the single largest lever (+0.176). Beyond it, gains come from *modeling and post-processing*, not more features — the static-statistic surrogates hit a hard ceiling on L2 (§6).

4. Temporal modeling & label alignment (Q3)

Alignment problem. Each file is a length-300, six-channel sequence with a *single* label. We align the sequential readings to that label in two complementary ways.

(a) Hand-crafted temporal features → GBDT. We summarize the 300-step sequence into descriptors that encode temporal dependence, so a gradient-boosted tree can use them: - **FFT band energies** — rhythmicity / dominant cadence; - **autocorrelation** (lags) — periodicity; - **sub-window pooling** — within-file phase/segment changes; - **jerk** (first difference of the mean channels) — motion transitions; - **zero-crossing rates** — oscillation frequency surrogate; - **catch22** — 22 canonical time-series characteristics (entropy, trend, distribution shape).

(b) Orientation “pseudo-gyroscope.” Because gravity is *not* removed and the device is wrist-worn, the per-second mean *x/y/z* trace the **wrist’s orientation trajectory** over the 300 seconds. The **time derivative of that orientation vector approximates the missing gyroscope** (angular velocity). We derive pitch/roll/inclination and their per-second changes (reversals, angular speed), which inject genuinely new rotational-dynamics signal for the hard L2 class.

(c) Sequence models. To capture temporal dependencies the summaries miss, we train models that consume the raw 6×300 directly: **CNN-BiLSTM** (1-D convolutions learn derivative-like and multi-scale kernels; the BiLSTM captures long-range dependence) and **InceptionTime** (multi-scale parallel convolutions). These contribute as ensemble members.

The feature ablation (§6) confirms this matters: **five of the six feature groups most critical to L2 are time-domain** (jerk, sub-window, FFT, zero-cross, autocorr).

5. Final pipeline

```
RAW (6 channels × 300 s, 1 Hz, gravity-laden, wrist)
├─ FEATURES: 271 stats + catch22 + 11 families + orientation pseudo-gyro dynamics
├─ BASE MODELS: P1 = combo_full_v2 (LightGBM stacker, 805 cols)
│               P2 = hier_v6 (Coarse → Fine_walk[LGBM+XGB] → Fine_other)
├─ BLEND:       0.842 · P1 + 0.158 · P2
├─ CALIBRATE:   per-class isotonic regression (5-fold GroupKFold OOF)
├─ PRIOR-CORR:  × (test_prior / train_prior)^2.0 [Saerens label-shift, label-free] ← test-time
├─ ORIENT:      inject orientation pseudo-gyro source into the L2 column (w = 0.15)
└─ THRESHOLD:   per-class log-weight ("robust" config) → argmax
                                                         → public LB 0.8234
```

The whole post-base pipeline is **deterministic and reproduces from frozen OOFs** (reproduce_final.py).

6. Ablation study (Q4)

6.1 Feature-group ablation (remove one group from the 271-feature catalog; GroupKFold-5)

Baseline `lgbm_full_v1` = **0.7091**; per-class L2 = 0.173. Noise floor (constant quality group) $\Rightarrow |\Delta| < 0.005$ is not meaningful.

Removed group	Δ macro-F1	Δ L2
jerk	-0.0088	-0.051
subwindow	-0.0068	-0.035
zerocross	-0.0060	-0.026
basic_stats	-0.0050	-0.057
fft	-0.0041	-0.034
magnitude	-0.0041	—
autocorr	-0.0031	-0.025
gravity	-0.0027	—
per_file_norm	-0.0010 (noise)	—

Finding: L2 depends on a *coalition* of time-domain groups that no single group replaces — which is why the static-feature L2 ceiling is ≈ 0.17 even with the full catalog (motivating the sequence models and the orientation pseudo-gyro). A counter-intuitive result: removing `basic_stats` *improved* L4 (+0.015) by cutting redundant high-intensity features the model was overfitting on its 142 files.

6.2 Pipeline-stage ablation (public LB — the core design choices)

Pipeline	OOF	Public LB	Δ LB
Blend + isotonic + threshold grid (base)	0.7880	0.8154	—
+ orientation pseudo-gyro	0.7873	0.8200	+0.0046
L2-injection			
+ test-prior correction (Saerens, $\beta=1$)	0.7808	0.8220	+0.0020
+ stronger test-prior correction ($\beta=2.0$)	0.7758	0.8234	+0.0014

*Threshold note: the base row uses the OOF-optimal “peak” thresholds (OOF 0.7880); from the orientation row onward we use the “robust” thresholds (OOF 0.7864) that generalize better on the public LB. So the orientation row’s OOF reflects the threshold switch **and** the injection (the injection alone is +0.0009 OOF on robust). The LB column is the ground-truth comparison.*

Key finding — OOF and LB diverge by design, not noise. The prior-correction stages have the *lowest* OOF yet the *highest* LB. The reason: the **test set has a different class prior than train** — Saerens-EM (label-free, estimated from the whole test set) puts L2 at ≈ 0.041 vs 0.033 in train and L3 at ≈ 0.073 vs 0.060. Our thresholds were tuned to the *train* prior, so train-OOF penalizes prior-adapted predictions while the actual test rewards them. Correcting toward the measured test prior — and amplifying it to $\beta=2.0$ (the EM estimate is conservative because the model is under-confident on the rare classes) — is the breakthrough from 0.8200 $\rightarrow$ 0.8234. Because the prior is estimated from the entire test set, the correction is expected to hold on the private split as well.

7. Results

Final standing: **public LB 0.8234** (sub_pc_b20.csv), up from a 0.8154 strong baseline. Progression and the role of each addition:

Submission	Public LB
sub_hier_v6_a842_grid_peak (GBDT stack + hierarchy + threshold grid)	0.8154
sub_robust_orient_inject_w15 (+ orientation pseudo-gyro)	0.8200
sub_robust_orient_L2_priorcorr (+ test-prior correction $\beta=1$)	0.8220
sub_pc_b20 (+ test-prior correction $\beta=2.0$)	0.8234

For reference, the majority/LR/LGBM baselines score 0.1331 / 0.5315 / 0.7091 (CV); the final pipeline is **+0.11 over the strongest provided baseline (0.7088)**.

8. What we explored that did not make the final pipeline

To be transparent about the search (and avoid clutter in the method): we additionally evaluated, under the same GroupKFold + nested-CV discipline, **domain-generalization / cross-user contrastive learning, GRU and evidential models, ROCKET/MultiRocket on orientation-rich channels, a from-scratch InceptionTime, and 2-D CNNs on Gramian-Angular-Field / spectrogram images**. Each carried *some* complementary L2 signal in isolation, but **none stacked**: every additional source competes for the same ~ 10 – 15 ambiguous L1 $\leftrightarrow$ L2 boundary samples, so the joint gain collapses to noise. This is consistent with a measured **Bayes-error floor** for L1/L2 at this 1 Hz / no-gyroscope resolution — i.e., the realizable accuracy ceiling is a property of the data, and the gains that *did* transfer came from new *physical* signal (orientation) and distribution adaptation (test-prior correction), not from more model capacity.

9. Reproducibility

- **Fast path (verifies 0.8234):** `python scripts/reproduce_final.py` or run notebooks/03_reproduce_final.ipynb — loads the committed frozen OOFs, runs the full final pipeline, asserts the winning class counts, and writes submissions/sub_pc_b20.csv.
- **Full path (from raw data):** README.md $\rightarrow$ *Full pipeline* documents every training stage (feature build $\rightarrow$ base models $\rightarrow$ hierarchical models $\rightarrow$ orientation source $\rightarrow$ final assembly).
- All randomness is seeded; CV is GroupKFold(5, groups=user_id) throughout.

Repository: <https://github.com/Ashurali/DM2026-Assignment-3-MKS>